

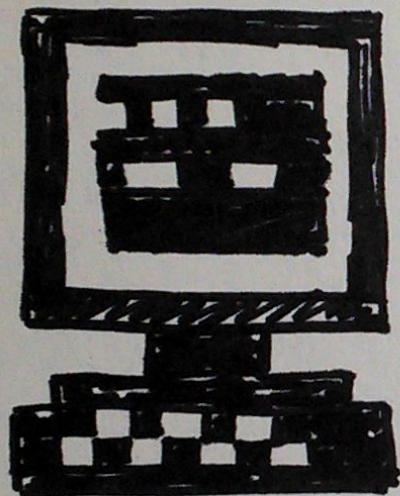
# PROGRAMMING EXPERIMENTS

# RULES OF PROGRAMMING EXPERIMENTS

1. it doesn't have to work
2. you don't have to finish it
3. you have to learn something

# the Recurse Center

New York!  
12 weeks!  
free!  
♡ ♡



the best  
programming  
community  
in the world

[recurse.com](http://recurse.com)

# EXPERIMENT 1: WRITE AN OPERATING SYSTEM

remember it doesn't have to work

~~Frontend  
Backend~~

~~Machine  
Learning~~

~~Using a  
database~~

~~JS~~

What's hard?

+

scary

~~Java~~

~~Web dev~~

~~Python~~

~~any new  
language~~

OPERATING

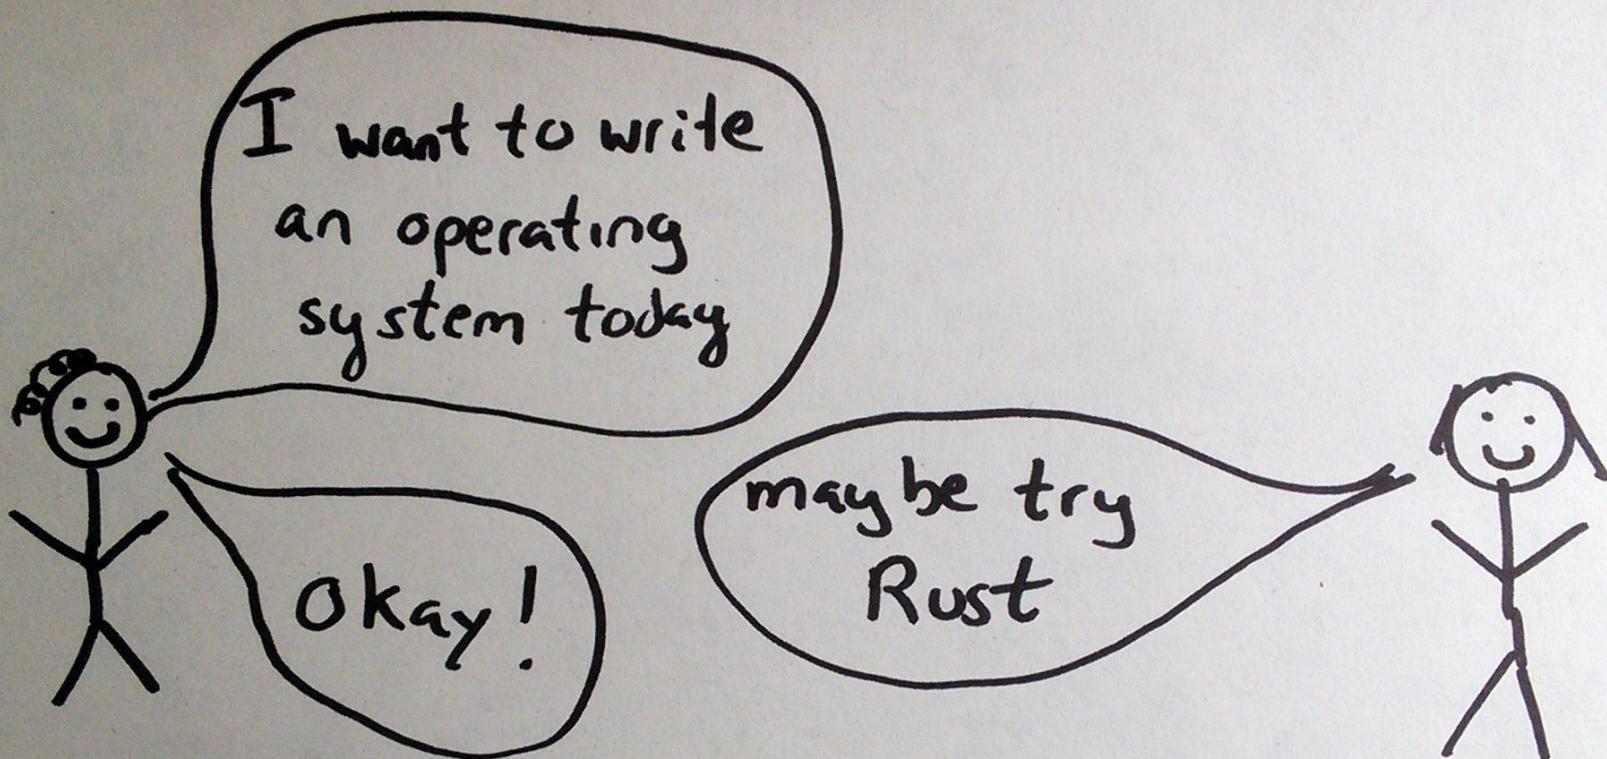
SYSTEMS

yay

amaze

!!

wow



me

Lindsey  
Kuper

Friday night at the  
Recurse Center

Rust lets you write  
your own

- malloc
  - threads
  - I/O
- etc  
etc



"~~freestanding mode~~"  
libcore

things you do not have when  
you write a kernel

- \* a hard drive
- \* files
- \* other processes
- \* a working keyboard
- \* memory protection
- \* any I/O
- \* libraries to use
- \* ANYTHING !!

"I'm just going to write  
a keyboard driver "

dramatic reading of  
"After 5 days, my OS doesn't  
crash when I press a key"



remember how we don't

have malloc? !! ! !

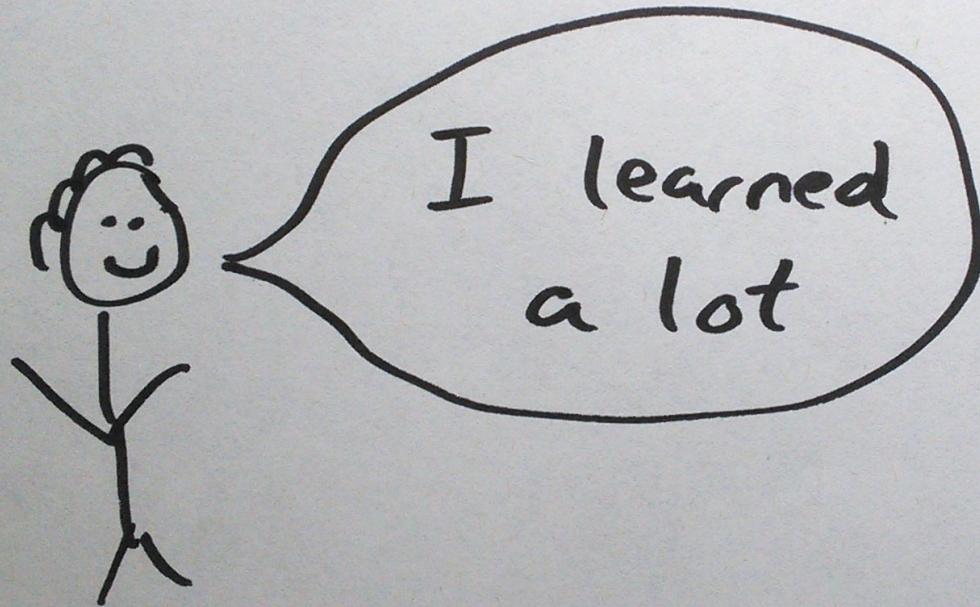
```
fn malloc(size: i32) {  
    return 42;  
}
```

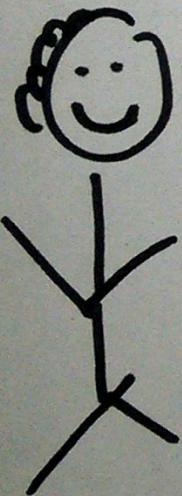
now we have our weird malloc

let  $x = \text{Box}::\text{new}(5)$   $\hookleftarrow^{x=5}$

let  $y = \text{Box}::\text{new}(88)$ ,  
now  
 $x = 88$  //

after 3 weeks





I'm trying to  
write a kernel  
and I have no  
idea what I'm  
doing also how  
do Rust strings  
work. help.

we can help!

#rust



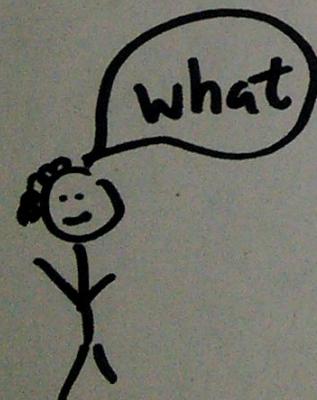
basically writing  
a toy kernel is

REALLY FUN

from: -@apple.com  
to: julia@jvns.ca  
subject: Opportunity with Apple

Hi,

I manage the Kernel Performance  
team at Apple...



# WHAT I LEARNED

- having more than once process is hard
- filesystems are hard
- operating systems are SO COMPLICATED

# EXPERIMENT 2: DATABASE SURGERY

# HOW DOES SQLITE WORK?

# FUN.SQLITE

id	word
1	greedy
2	greediness
3	greener

```
$ hexdump fun.sqlite
|.....{.n|
|.a.R.D.4.%....|
|.....|
|...y.n._.N.>.,.$|
|.....|
|.....F.|
|..EAcevedo.E...D|
|Accra's.D...CAcc|
|ra.C..#BAccentur|
|e's.B...AAccentu|
|re.A..!@Acapulco|
|'s.@...?Acapulco|
```

# A DATABASE IS A *TREE*

```
static MemPage *btreePageFromDbPage(DbPage *pDbPage, Pgno pgno, BtSha
    // actual code
    printf("Read a btree page, page number %d\n", pgno); // added by me
    // actual code
}
```

```
sqlite> select * from fun where id = 1;  
Read a btree page, page number 1  
Read a btree page, page number 5  
Read a btree page, page number 828  
Read a btree page, page number 10  
Read a btree page, page number 2  
Read a btree page, page number 76  
Read a btree page, page number 6  
1|A's
```

```
sqlite> select * from fun where id = 20;  
Read a btree page, page number 1  
Read a btree page, page number 5  
Read a btree page, page number 828  
Read a btree page, page number 10  
Read a btree page, page number 2  
Read a btree page, page number 76  
Read a btree page, page number 6  
20|Aaliyah
```

```
sqlite> select * from fun where id = 80000;
Read a btree page, page number 1
Read a btree page, page number 5
Read a btree page, page number 1198
Read a btree page, page number 992
Read a btree page, page number 2
Read a btree page, page number 1813
Read a btree page, page number 449
80000|scarfs
```

# WHAT I LEARNED

- databases tables are trees
- databases are made of pages
- i can read some of the SQLite source code!

# EXPERIMENT 3: WRITE A TCP STACK

# **EXPERIMENT 3: WRITE A TCP STACK IN PYTHON**



David Wolever and 1 other Retweeted



**Gary Bernhardt** @garybernhardt · 15 Jul 2015

Your **homework** is to implement UDP in Python using raw sockets (SOCK\_RAW) and get another machine to talk to you. Hard mode: implement **TCP**.



18



88

• • •

```
ip_header = IP(dst=dest_ip, src=src_ip)
syn = TCP(dport=80, sport=59333,
          ack=0, flags="S")
# Send the SYN packet to Google
response = srp(ip_header + syn)
```

# WHAT I LEARNED

- how TCP packets are put together!
- you can write a 10% working TCP from scratch in 2 weeks
- python can't keep up

# EXPERIMENT 4: CONCURRENCY

```
int counter;

void *AddThings(void *threadid) {
    for (int i = 0; i < 10000; i++)
        counter += 1;
    pthread_exit(NULL);
}
```

# WRONG ANSWER

# MUTEX

```
pthread_mutex_lock(&mutex);  
counter += 1;
```

# 'ATOM'

```
__sync_add_and_fetch(&counter, 1);
```

# WHAT I LEARNED

- atoms are faster than mutexes

"can you discuss the pros and cons of using a lock-free approach  
for implementing a thread-safe hashmap?"

# I BLOG MY EXPERIMENTS

## Diving into HDFS

in [strace](#)

---

Yesterday I wanted to start learning about how HDFS (the Hadoop Distributed File System) works internally. I knew that

- It's distributed, so one file may be stored across many different machines
- There's a *namenode*, which keeps track of where all the files are stored
- There are *data nodes*, which contain the actual file data

But I wasn't quite sure how to get started! I knew how to navigate the filesystem from the command line (`hadoop fs -ls /`, and friends), but not how to figure out how it works internally.

**DO ENOUGH  
EXPERIMENTS**

**END UP WITH ACTUAL  
KNOWLEDGE**